



Un viaje desde un Single Point of Failure hacia una solución de alta disponibilidad usando InnoDB ClusterSet

Francisco Bordenave





Agenda

Intro

El problema

Requerimientos

Desafíos

Solución propuesta

Procedimiento

Consideraciones finales

Preguntas?





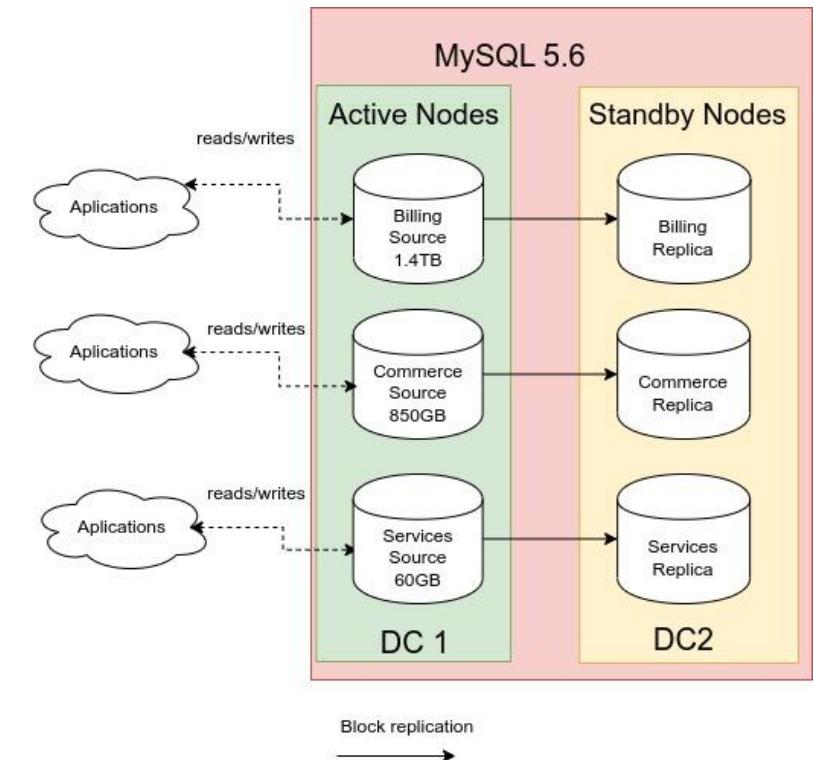
El problema



El problema (basado en una historia real)

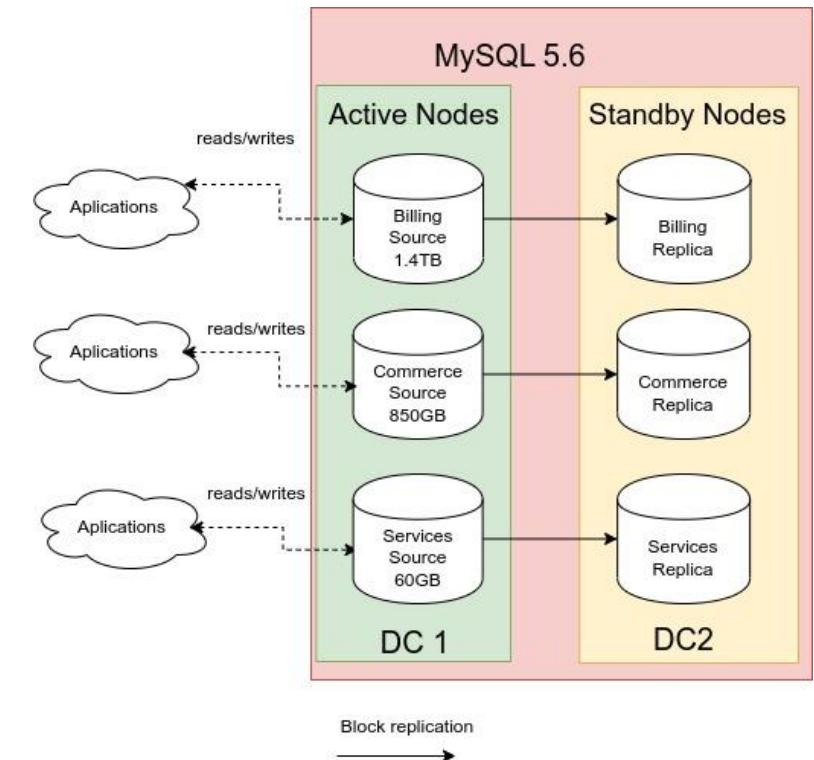
Escenario

MySQL 5.6 EoL por años



Escenario

**MySQL 5.6 EoL por años
Single point of failure**

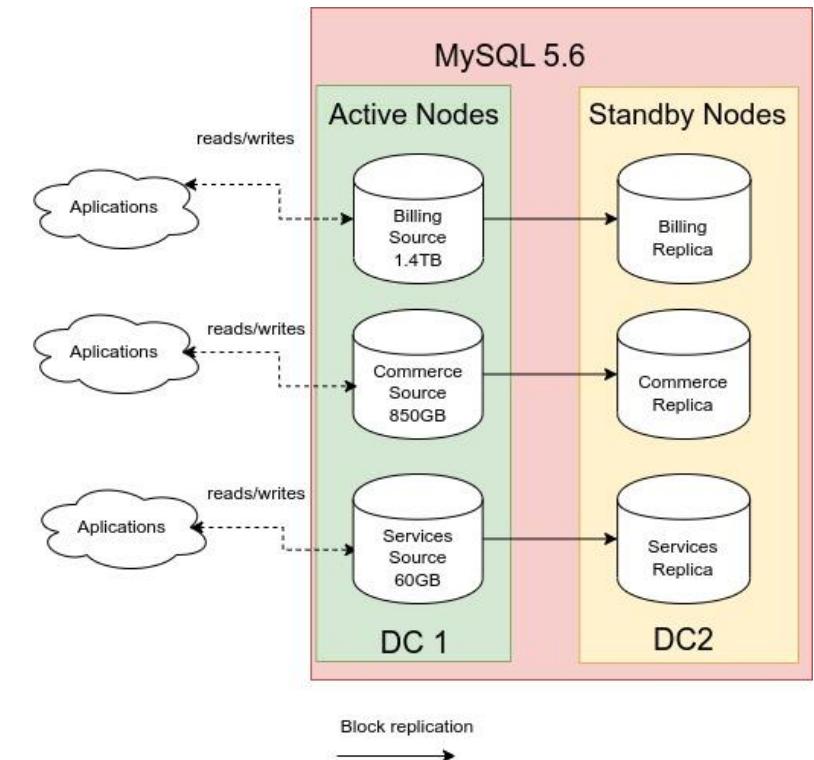


Escenario

MySQL 5.6 EoL por años

Single point of failure

Failover manual



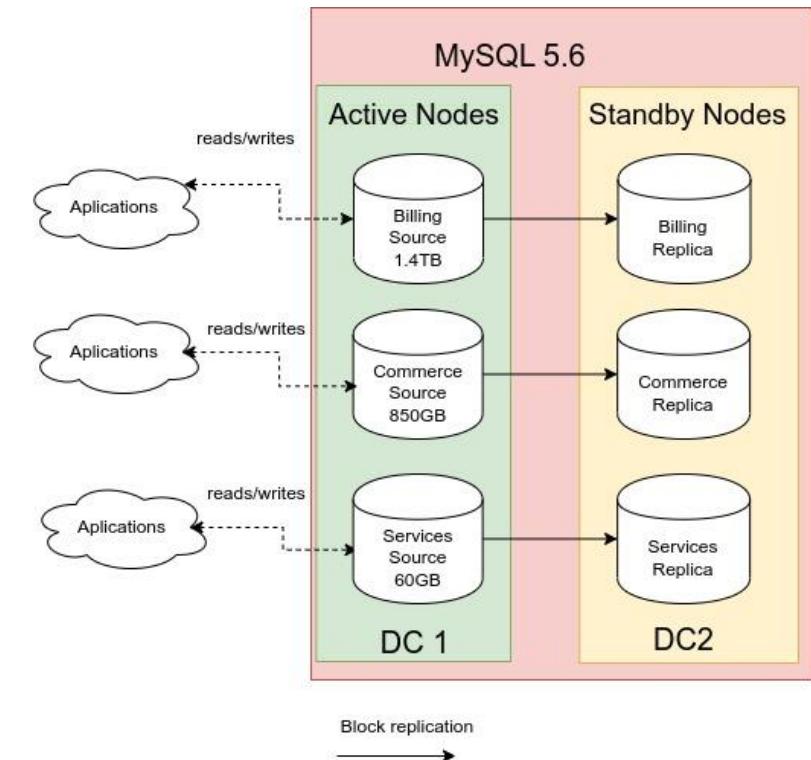
Escenario

MySQL 5.6 EoL por años

Single point of failure

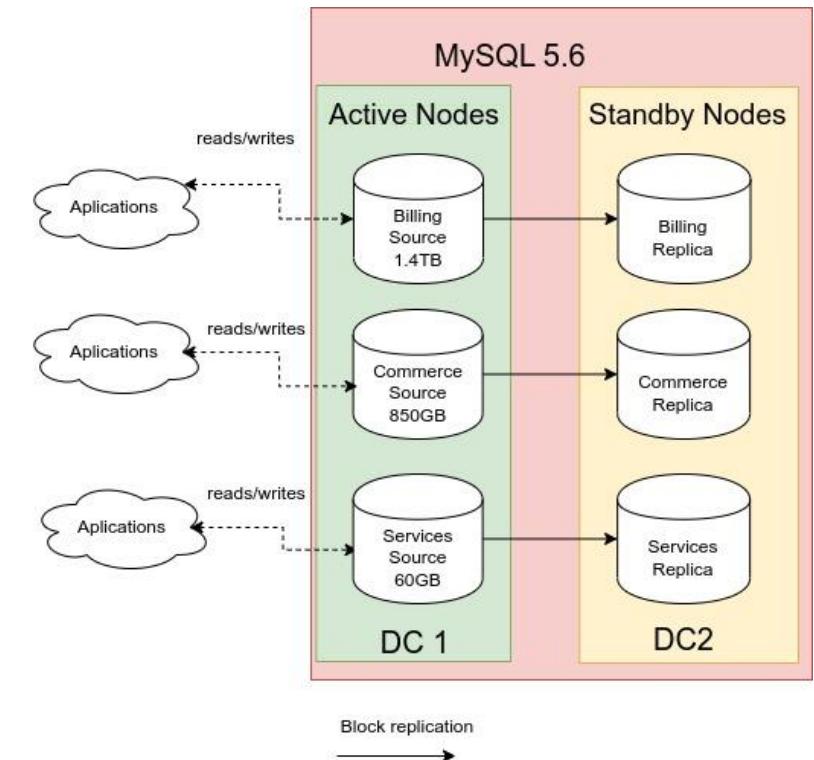
Failover manual

Sin balanceo de carga



Escenario

MySQL 5.6 EoL por años
Single point of failure
Failover manual
Sin balanceo de carga
Administración compleja



Escenario

MySQL 5.6 EoL por años

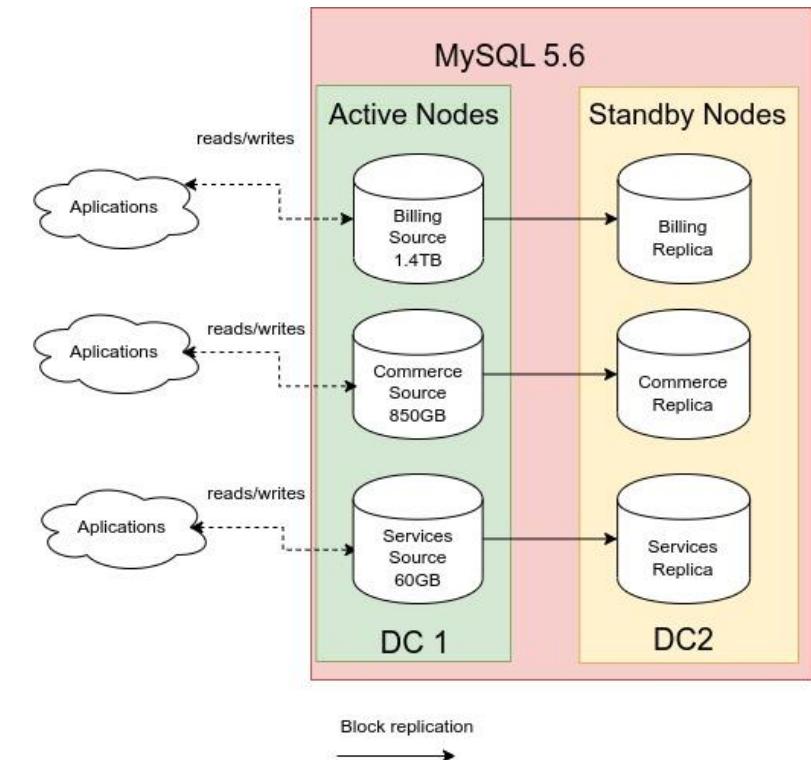
Single point of failure

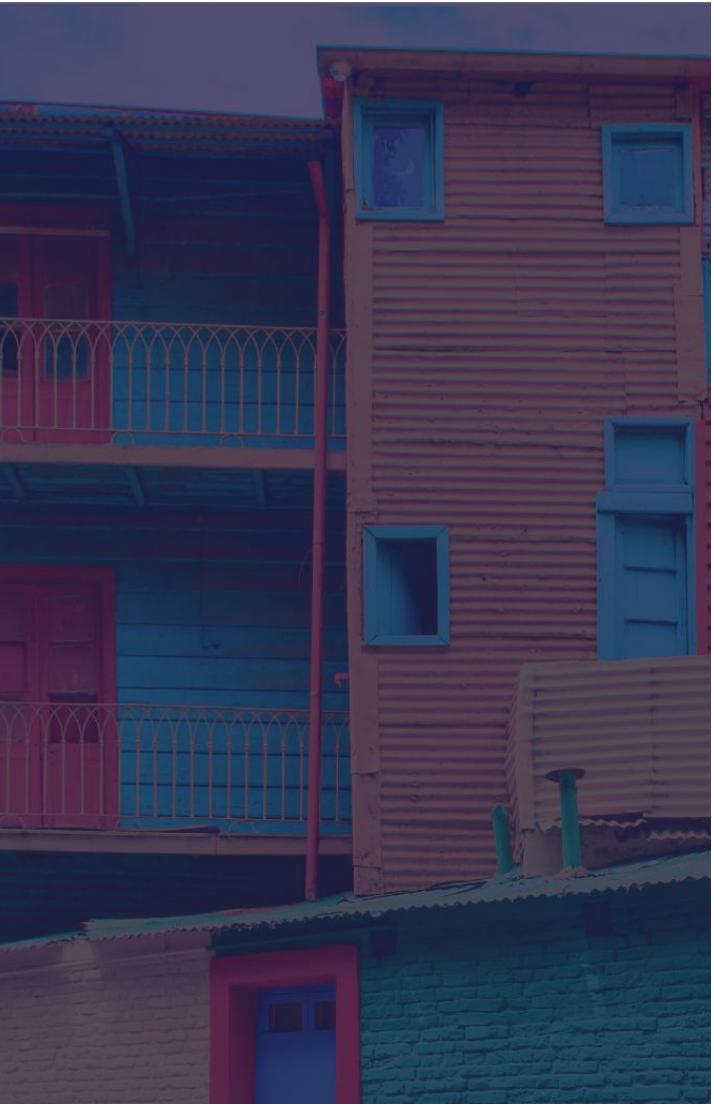
Failover manual

Sin balanceo de carga

Administración compleja

Mal manejo de DR





Requerimientos



Requerimientos

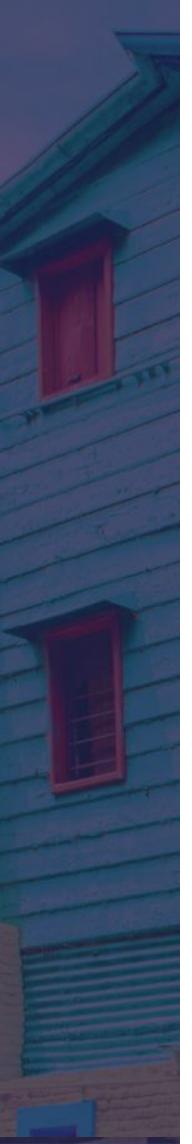
Actualizar a 8.0 - SLA



Requerimientos

Actualizar a 8.0 - SLA

Alta disponibilidad



Requerimientos

Actualizar a 8.0 - SLA

Alta disponibilidad

Tolerante a fallos de DC - Planificar DR



Requerimientos

Actualizar a 8.0 - SLA

Alta disponibilidad

Tolerante a fallos de DC - Planificar DR

0 Downtime durante el upgrade



Requerimientos

Actualizar a 8.0 - SLA

Alta disponibilidad

Tolerante a fallos de DC - Planificar DR

0 Downtime durante el upgrade

Fácil mantenimiento



Requerimientos

Actualizar a 8.0 - SLA

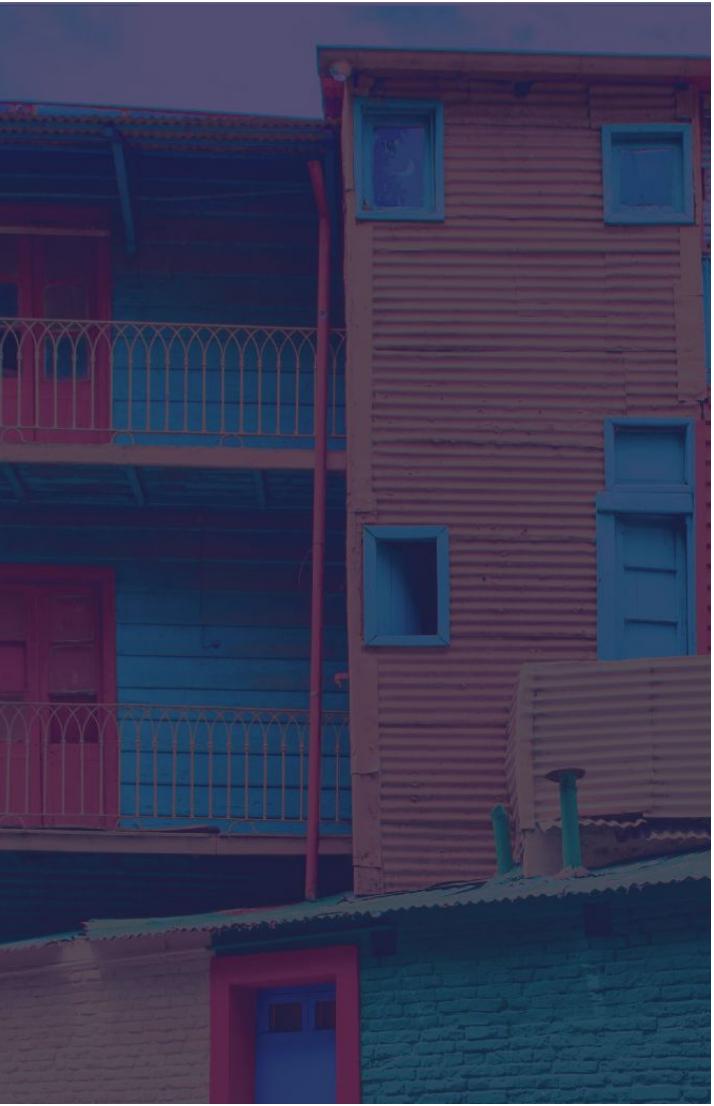
Alta disponibilidad

Tolerante a fallos de DC - Planificar DR

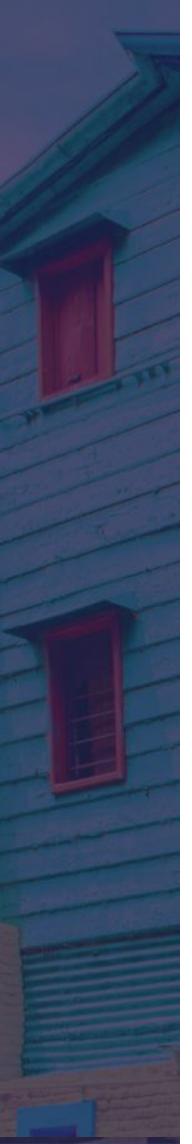
0 Downtime durante el upgrade

Fácil mantenimiento

.... y más



Desafíos



Desafíos

Actualización intermedia a 5.7



Desafíos

Actualización intermedia a 5.7

Consolidar servicios para reducir complejidad



Desafíos

Actualización intermedia a 5.7

Consolidar servicios para reducir complejidad

Asegurar compatibilidad



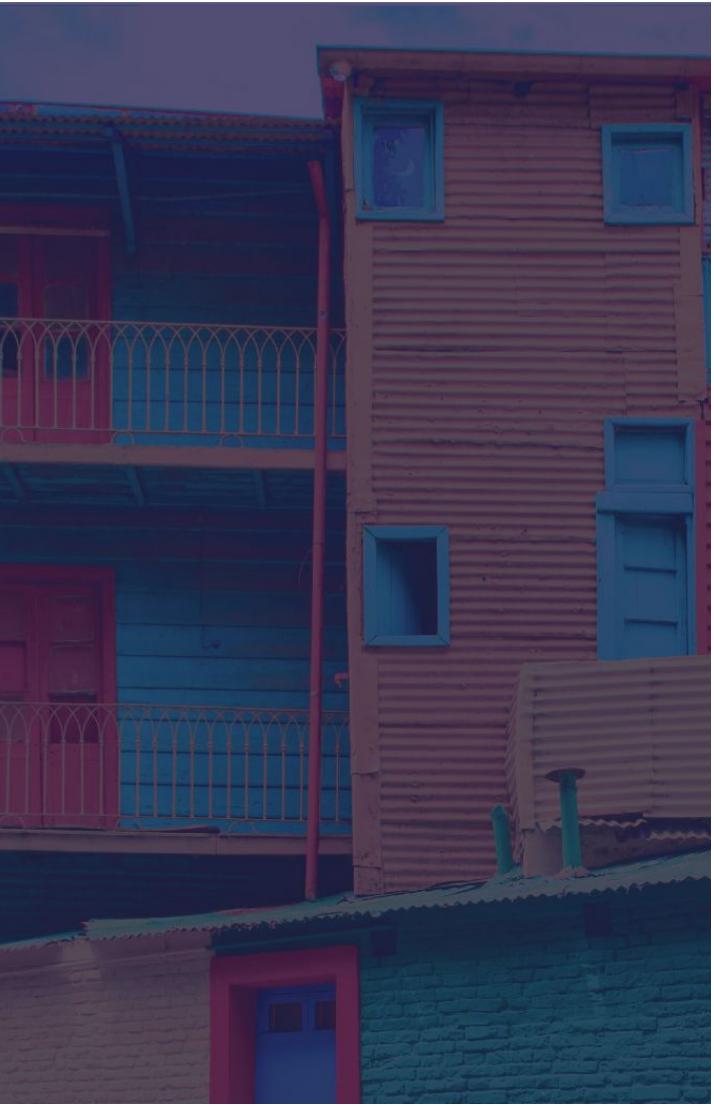
Desafíos

Actualización intermedia a 5.7

Consolidar servicios para reducir complejidad

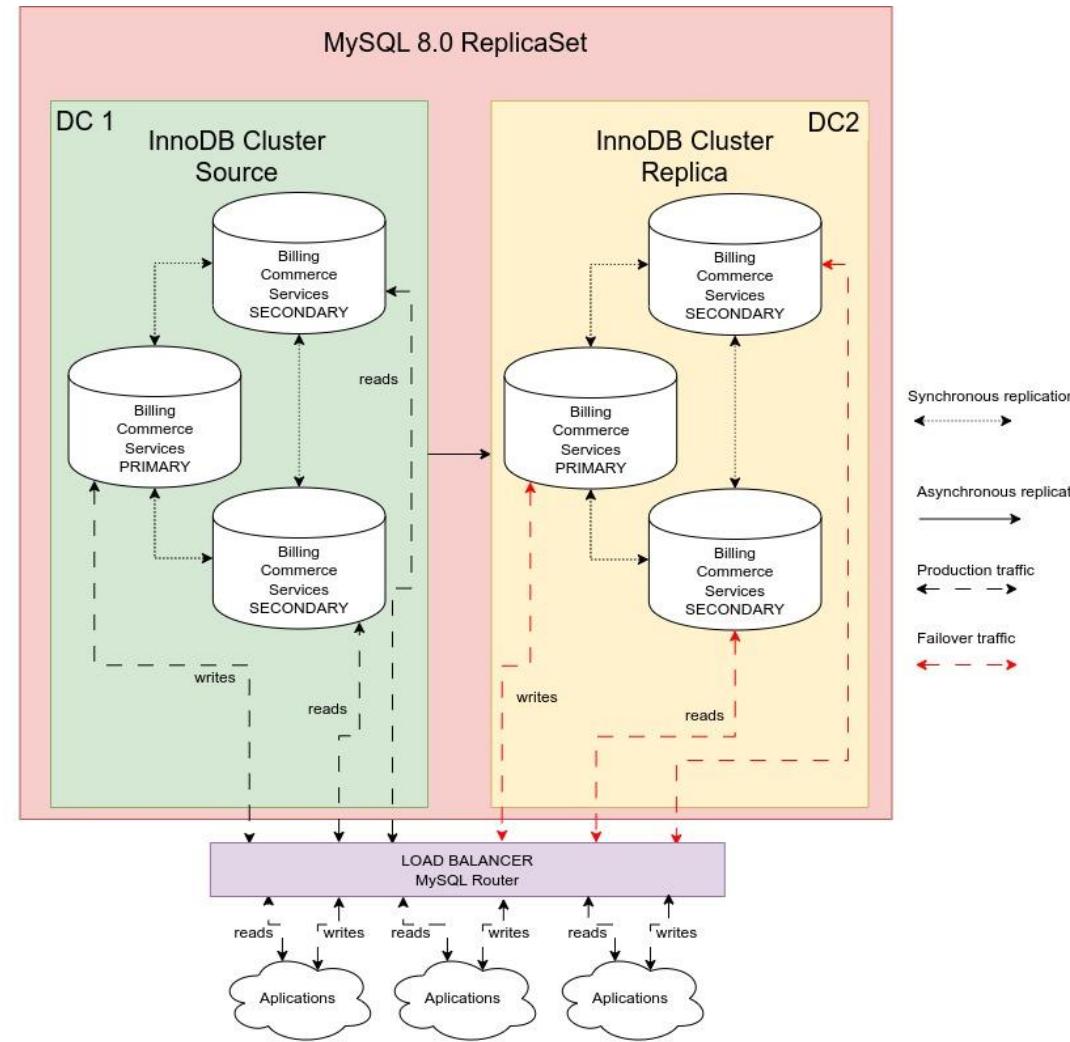
Asegurar compatibilidad

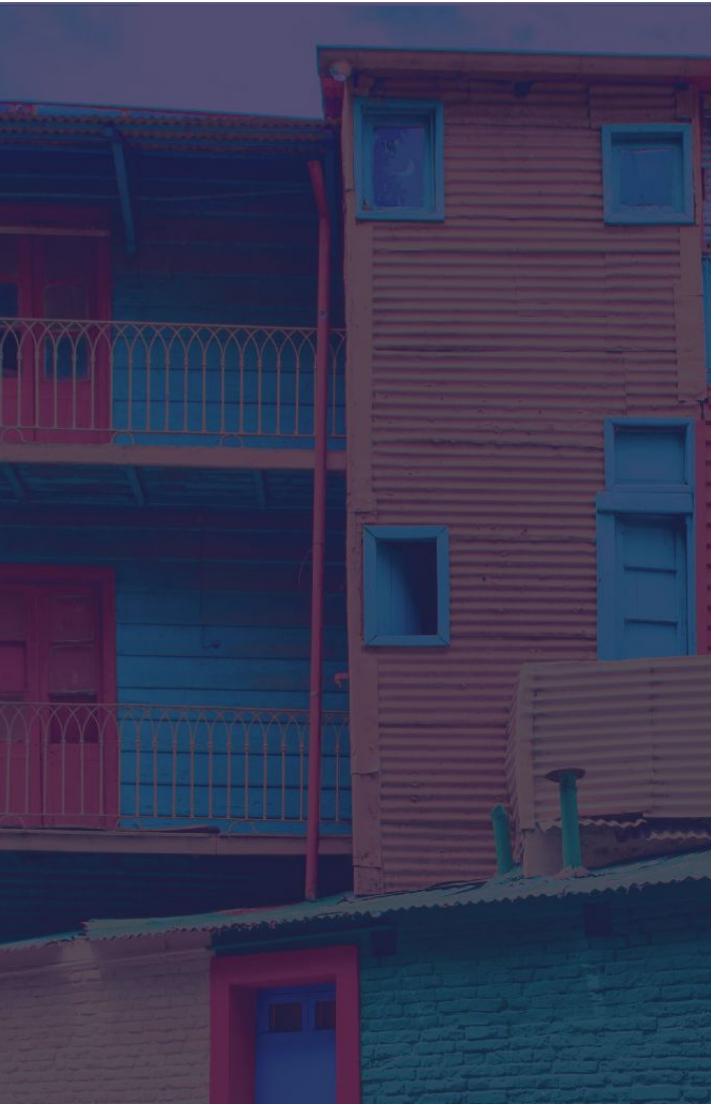
Mantener consistencia de datos durante la migración



Solución propuesta

Innodb ClusterSet



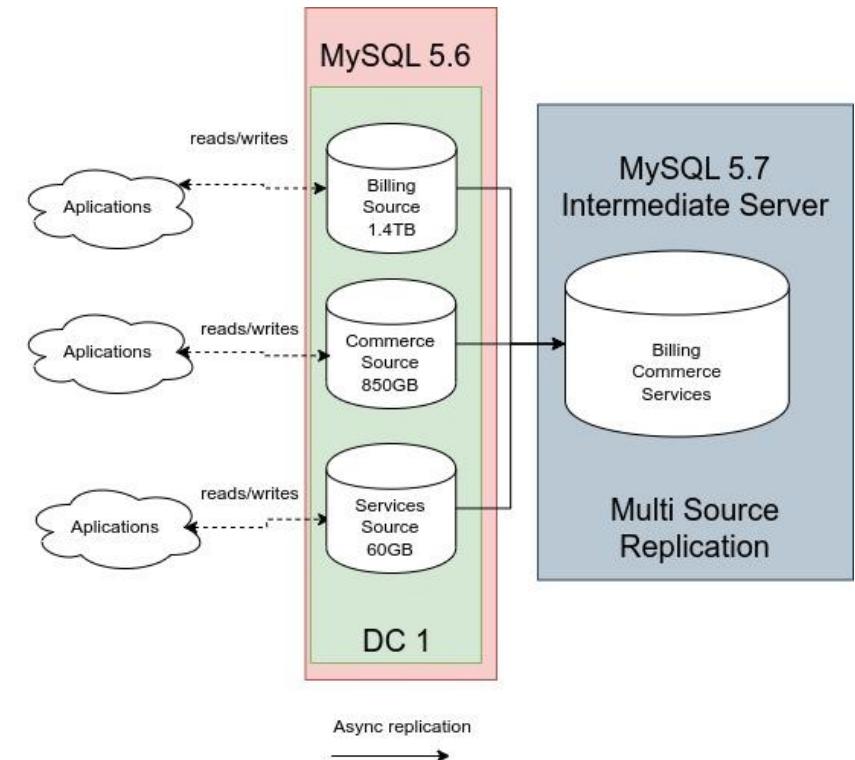


Procedimiento

Procedimiento

Crear un servidor intermedio

Testear upgrade a 5.7

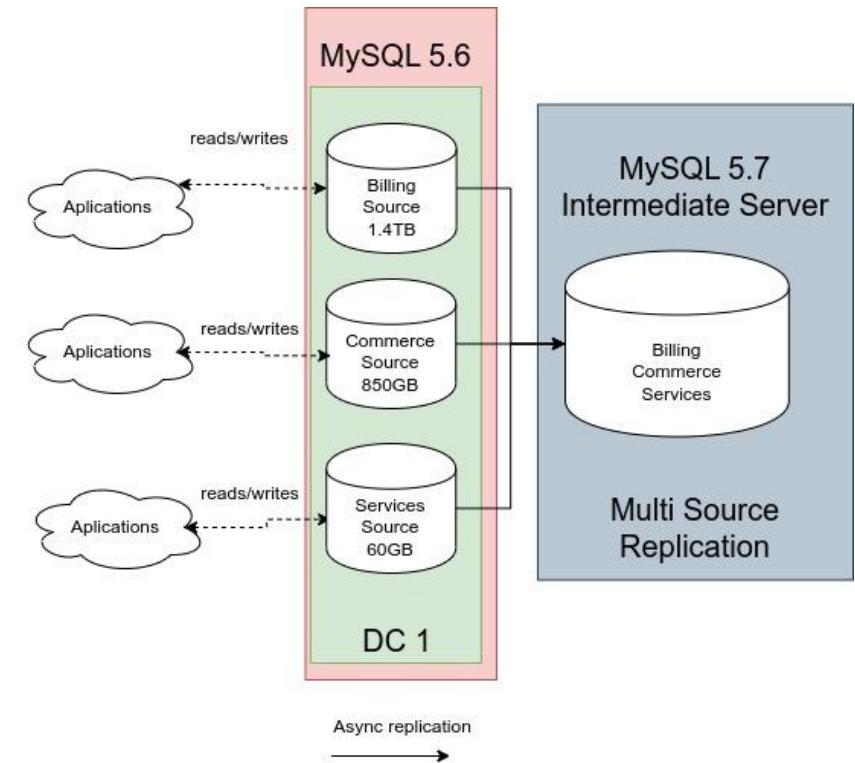


Procedimiento

Crear un servidor intermedio

Testear upgrade a 5.7

Testear compatibilidad con 5.6



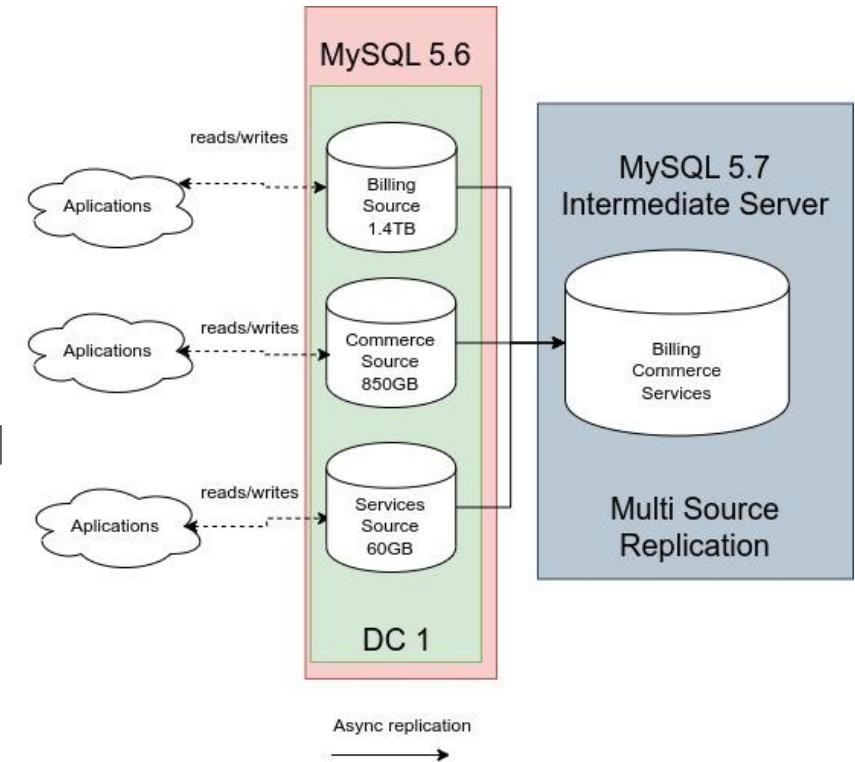
Procedimiento

Crear un servidor intermedio

Testear upgrade a 5.7

Testear compatibilidad con 5.6

Consolidar bases de datos para reducir complejidad



Procedimiento

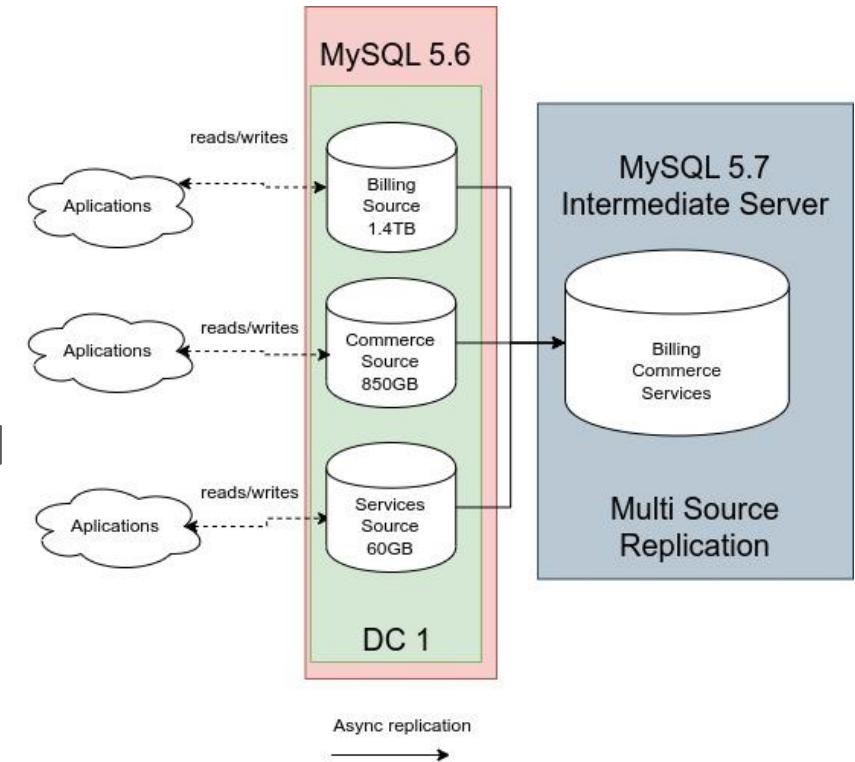
Crear un servidor intermedio

Testear upgrade a 5.7

Testear compatibilidad con 5.6

Consolidar bases de datos para reducir complejidad

No se pueden usar los secundarios como backup



Procedimiento

Crear un servidor intermedio

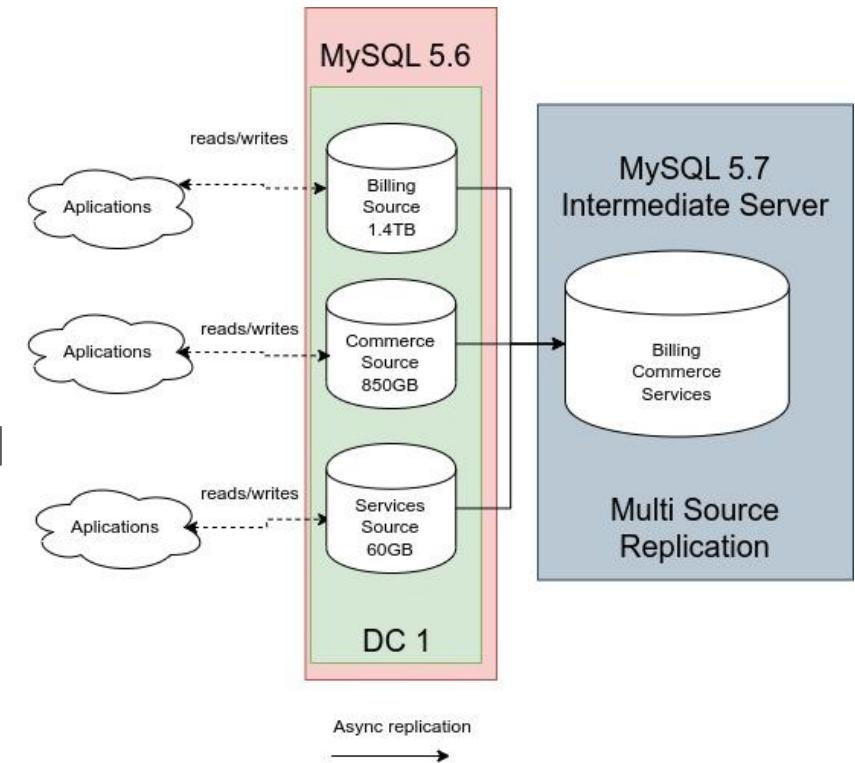
Testear upgrade a 5.7

Testear compatibilidad con 5.6

Consolidar bases de datos para reducir complejidad

No se pueden usar los secundarios como backup

Usar XtraBackup para la DB más grande



Procedimiento

Crear un servidor intermedio

Testear upgrade a 5.7

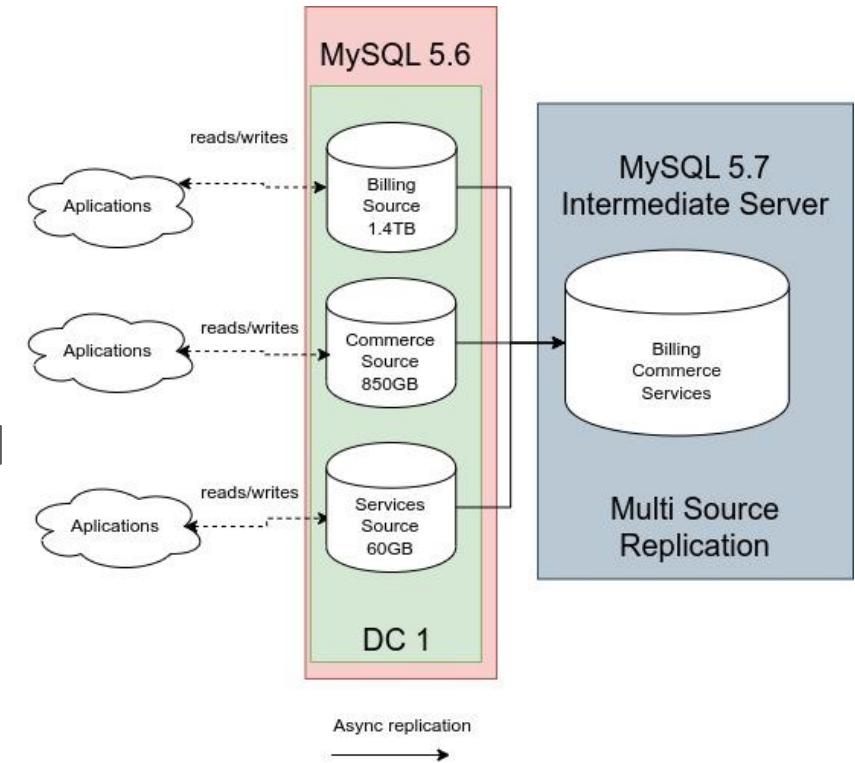
Testear compatibilidad con 5.6

Consolidar bases de datos para reducir complejidad

No se pueden usar los secundarios como backup

Usar XtraBackup para la DB más grande

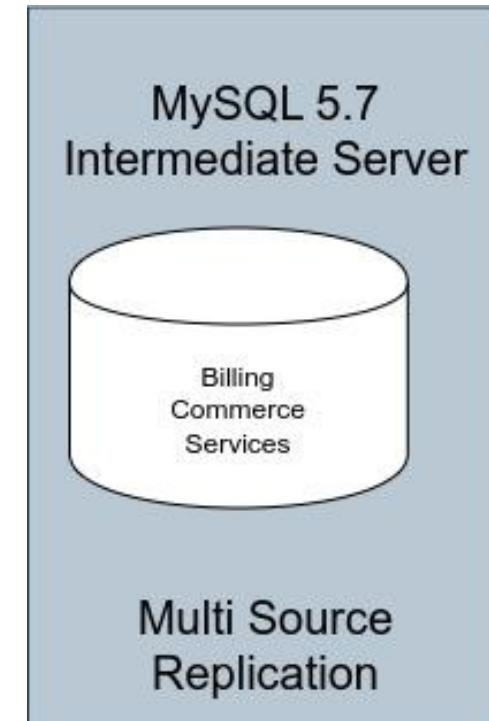
Usar mydumper/myloader para las DBs restantes



Procedimiento

Creación del servidor intermedio

```
my.cnf  
log-bin  
log-slave-updates  
binlog-format=row
```



Procedimiento

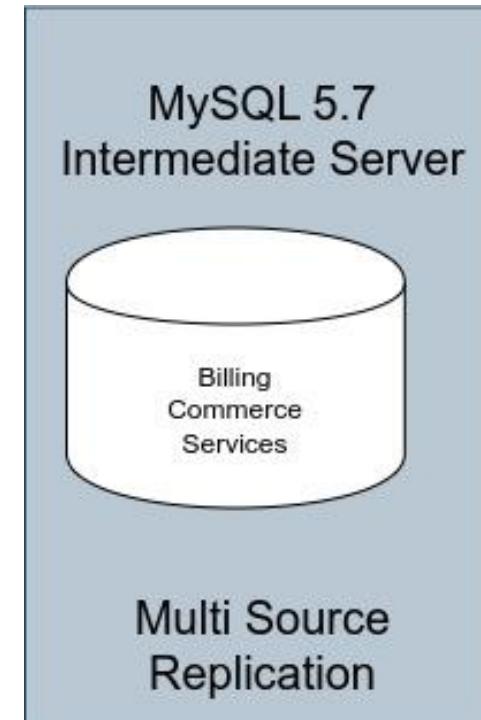
Creación del servidor intermedio

Origen - billing

```
xtrabackup --user='bkpuser' --password='s3cret' --backup  
--target-dir=./ --stream=xbstream | nc 10.50.30.250 3306
```

Servidor intermedio

```
nc -l 9999 xbstream -x -C /var/lib/mysql/  
xtrabackup --prepare --target-dir=/data/backups/  
mysql_upgrade -s
```



Procedimiento

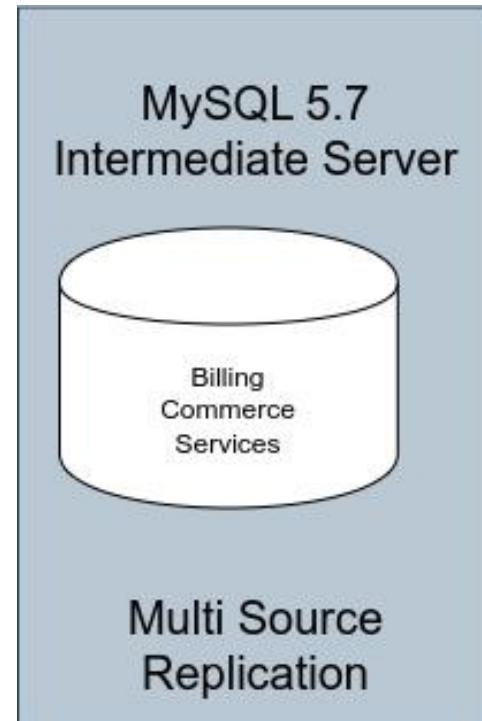
Creación del servidor intermedio

Origen - resto de las DBs

```
mydumper -v 2 -t 4 -G -E -R -o /var/lib/mysql/backup --regex  
'^^(audit_service_production\.|backend_service_production\.|bco  
mmerce_service_production\.)'  
  
mydumper -v 2 -t 4 -G -E -R -o /var/lib/mysql/backup --regex  
'^^(loyalty_service_production\.|support_service_production\.|v  
pos_service_production\.)'
```

Servidor intermedio

```
myloader --host=127.0.0.1 --directory=/var/lib/mysql/backup  
--threads=8
```



Procedimiento

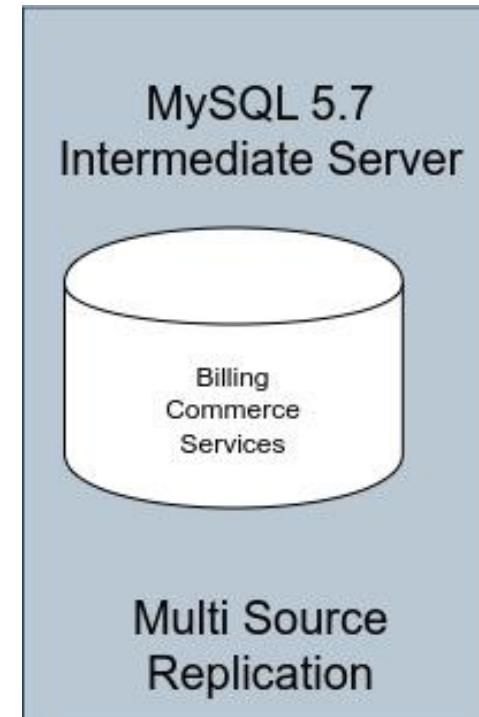
Creación del servidor intermedio

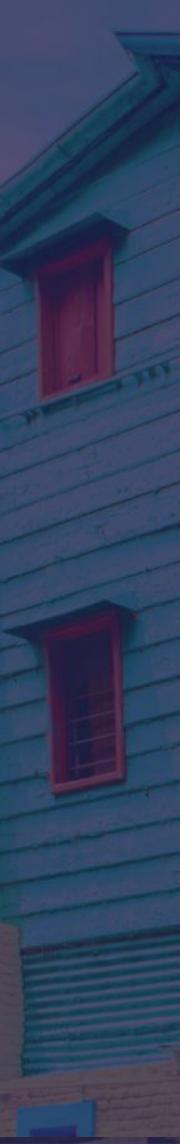
Servidor intermedio

```
CHANGE MASTER TO master_host='10.50.30.11',
master_user='replication', master_password='r3pl1c4t10n',
master_log_file='mysql-bin.002546', master_log_pos=570428664
for channel 'commerce';
```

```
CHANGE MASTER TO master_host='10.50.30.21',
master_user='replication', master_password='r3pl1c4t10n',
master_log_file='mysql-bin.000152', master_log_pos=36271765
for channel billing;
```

```
CHANGE MASTER TO master_host='10.50.30.7',
master_user='replication', master_password='r3pl1c4t10n',
master_log_file='mysql-bin.001611', master_log_pos=936303055
for channel 'services';
```





Procedimiento

El servidor 5.7 está listo - Y ahora?



Procedimiento

El servidor 5.7 está listo - Y ahora?

Nuevos desafíos



Procedimiento

El servidor 5.7 está listo - Y ahora?

Nuevos desafíos

Testear compatibilidad con 8.0

Procedimiento

El servidor 5.7 está listo - Y ahora?

Nuevos desafíos

Testear compatibilidad con 8.0

```
mysqlsh -- util checkForServerUpgrade
```

Procedimiento

El servidor 5.7 está listo - Y ahora?

Nuevos desafíos

Testear compatibilidad con 8.0

```
mysqlsh -- util checkForServerUpgrade
```

Group Replication requiere GTID

Procedimiento

El servidor 5.7 está listo - Y ahora?

Nuevos desafíos

Testear compatibilidad con 8.0

Group Replication requiere GTID

```
mysqlsh -- util checkForServerUpgrade  
ASSIGN_GTIDS_TO_ANONYMOUS_TRANSACTIONS
```

Preparando la instancia

```
mysqlsh -- util checkForServerUpgrade user@localhost:3306 --output-format=JSON
The MySQL server at example.com:3306, version
5.7.33-enterprise-commercial-advanced - MySQL Enterprise Server - Advanced Edition (Commercial),
will now be checked for compatibility issues for upgrade to MySQL 8.0.29...
1) Usage of old temporal type
   No issues found
2) Usage of db objects with names conflicting with new reserved keywords
   Warning: The following objects have names that conflict with new reserved keywords.
   Ensure queries sent by your applications use `quotes` when referring to them or they will result in errors.
   More information: https://dev.mysql.com/doc/refman/en/keywords.html
   dbtest.System - Table name
   dbtest.System.JSON_TABLE - Column name
   dbtest.System.cube - Column name
```

Configuración del server 8.0

```
[mysqld]
bind-address = 0.0.0.0
report-host = 10.50.30.200 #change this value on each 8.0 instance
server_id = 1001 #change this value on each 8.0 instance
gtid_mode = ON
enforce_gtid_consistency = ON
log-replica-updates = ON
group_replication_start_on_boot=OFF
binlog-checksum = NONE
binlog_transaction_dependency_tracking = WRITESET
replica_preserve_commit_order = 1
replica_parallel_type = LOGICAL_CLOCK
transaction-write-set-extraction = XXHASH64
```

Creando el InnoDB ClusterSet - DC1

```
CREATE USER 'clusteradm'@'10.%';
GRANT ALL PRIVILEGES ON *.* TO 'clusteradm'@'10.%';
GRANT CLONE_ADMIN, CONNECTION_ADMIN, CREATE USER, EXECUTE, FILE, GROUP_REPLICATION_ADMIN,
PERSIST_RO_VARIABLES_ADMIN, PROCESS, RELOAD, REPLICATION_CLIENT, REPLICATION_SLAVE,
REPLICATION_APPLIER, REPLICATION_SLAVE_ADMIN, ROLE_ADMIN, SELECT, SHUTDOWN,
SYSTEM_VARIABLES_ADMIN ON *.* TO 'clusteradm'@'10.%' WITH GRANT OPTION;
```

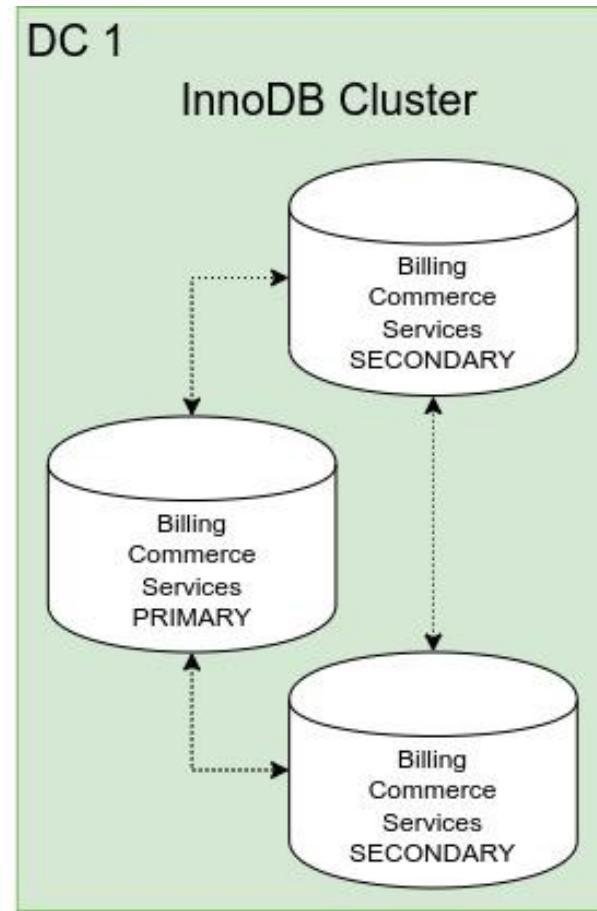
Creando el InnoDB ClusterSet - DC1

```
mysqlsh -uclusteradm -h10.50.30.101 -ppass  
dba.configureInstance()  
  
Configuring local MySQL instance listening at port 3306 for use in an InnoDB cluster...  
  
This instance reports its own address as 10.50.30.101:3306  
applierWorkerThreads will be set to the default value of 4.  
  
The instance '10.50.30.101:3306' is valid to be used in an InnoDB cluster.  
  
The instance '10.50.30.101:3306' is already ready to be used in an InnoDB cluster.  
  
Successfully enabled parallel appliers.  
  
dba.checkInstanceConfiguration()  
  
Validating local MySQL instance listening at port 3306 for use in an InnoDB cluster...  
  
...  
  
The instance '10.50.30.101:3306' is valid to be used in an InnoDB cluster.  
  
{  
    "status": "ok"
```

Creando el InnoDB ClusterSet - DC1

```
mysqlsh -uclusteradm -h10.50.30.101 -ppass
var cluster = dba.createCluster('GR_DC1');
cluster.addInstance({user:'clusteradm', host:'10.50.30.103', password:'pass'});
cluster.addInstance('clusteradm@10.50.30.105');
cluster.status()
{
    "clusterName": "GR_DC1",
    "defaultReplicaSet": {
        "name": "default",
        "primary": "10.50.30.101:3306",
        "status": "OK",
        "statusText": "Cluster is ONLINE and can tolerate up to ONE failure.",
        "topology": {
            "10.50.30.101:3306": {
                "address": "10.50.30.101:3306",
                "memberRole": "PRIMARY",
                "mode": "R/W",
                ...
            },
            "10.50.30.103:3306": {
                "address": "10.50.30.103:3306",
                "memberRole": "SECONDARY",
                ...
            }
        }
    }
}
```

Creando el InnoDB ClusterSet - DC1

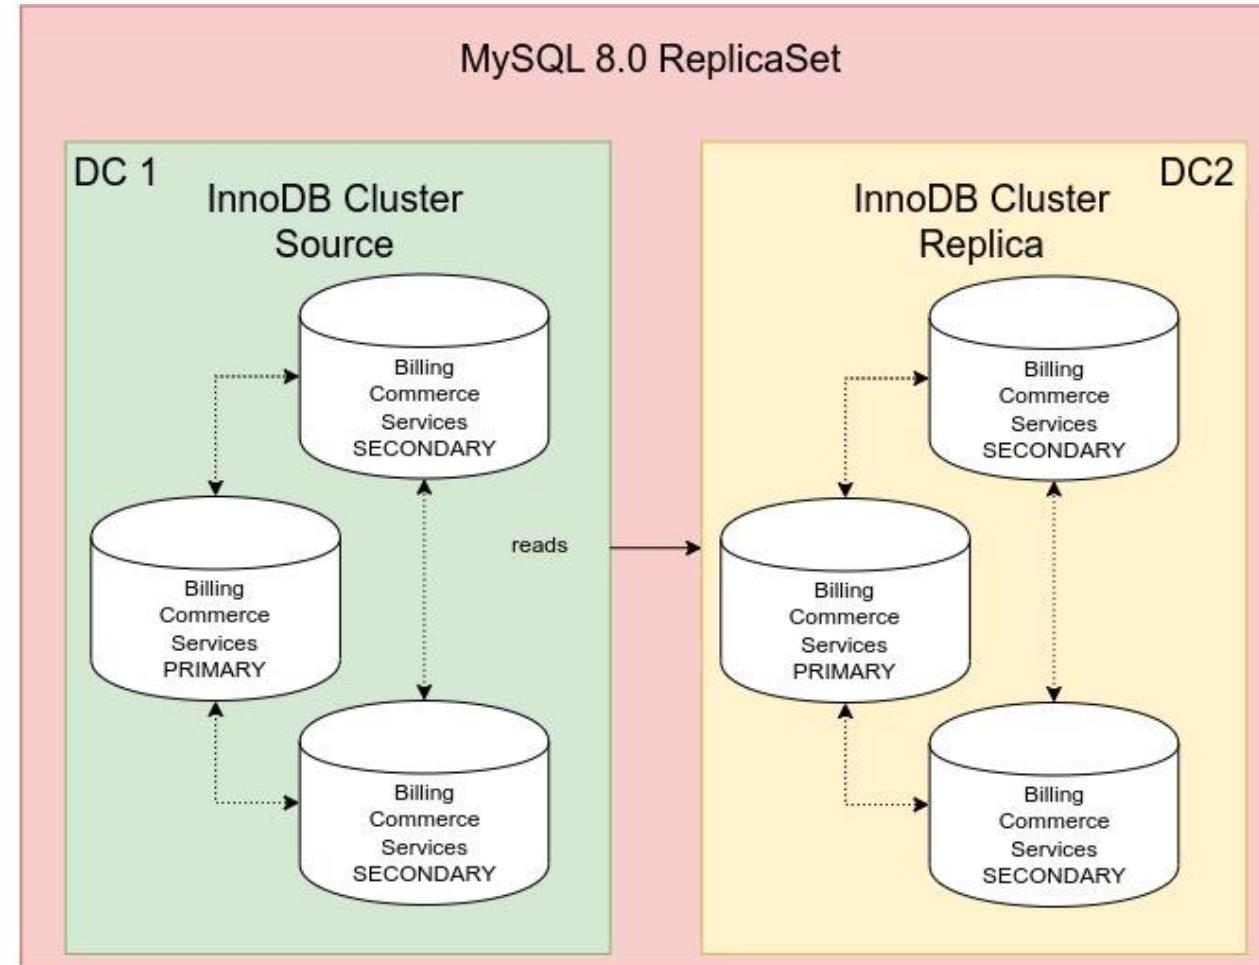


Creando el InnoDB ClusterSet - DC2 (Replica Cluster)

```
mysqlsh> var clusterSet = cluster.createClusterSet("GR_DC2");
mysqlsh> var clusterDC2 = clusterSet.createReplicaCluster("clusteradm@10.50.30.102", "GR_DC2");
mysqlsh> clusterDC2.addInstance('clusteradm@10.50.30.112');
mysqlsh> clusterDC2.addInstance('clusteradm@10.50.30.114');
mysqlsh> clusterDC2.status()

mysqlsh> clusterSet.status();
{
  "clusters": {
    "GR_DC1": {
      "clusterRole": "PRIMARY",
      "globalStatus": "OK",
      "primary": "10.50.30.101:3306"
    },
    "GR_DC2": {
      "clusterRole": "REPLICA",
      "clusterSetReplicationStatus": "OK",
      "globalStatus": "OK"
    }
  },
  "domainName": "GR_DC2",
  "globalPrimaryInstance": "10.50.30.101:3306",
  "primaryCluster": "GR_DC1",
  "status": "HEALTHY",
  "statusText": "All Clusters available."
}
```

Creando el InnoDB ClusterSet - DC2 (Replica Cluster)



Creando el InnoDB ClusterSet - MySQL Router

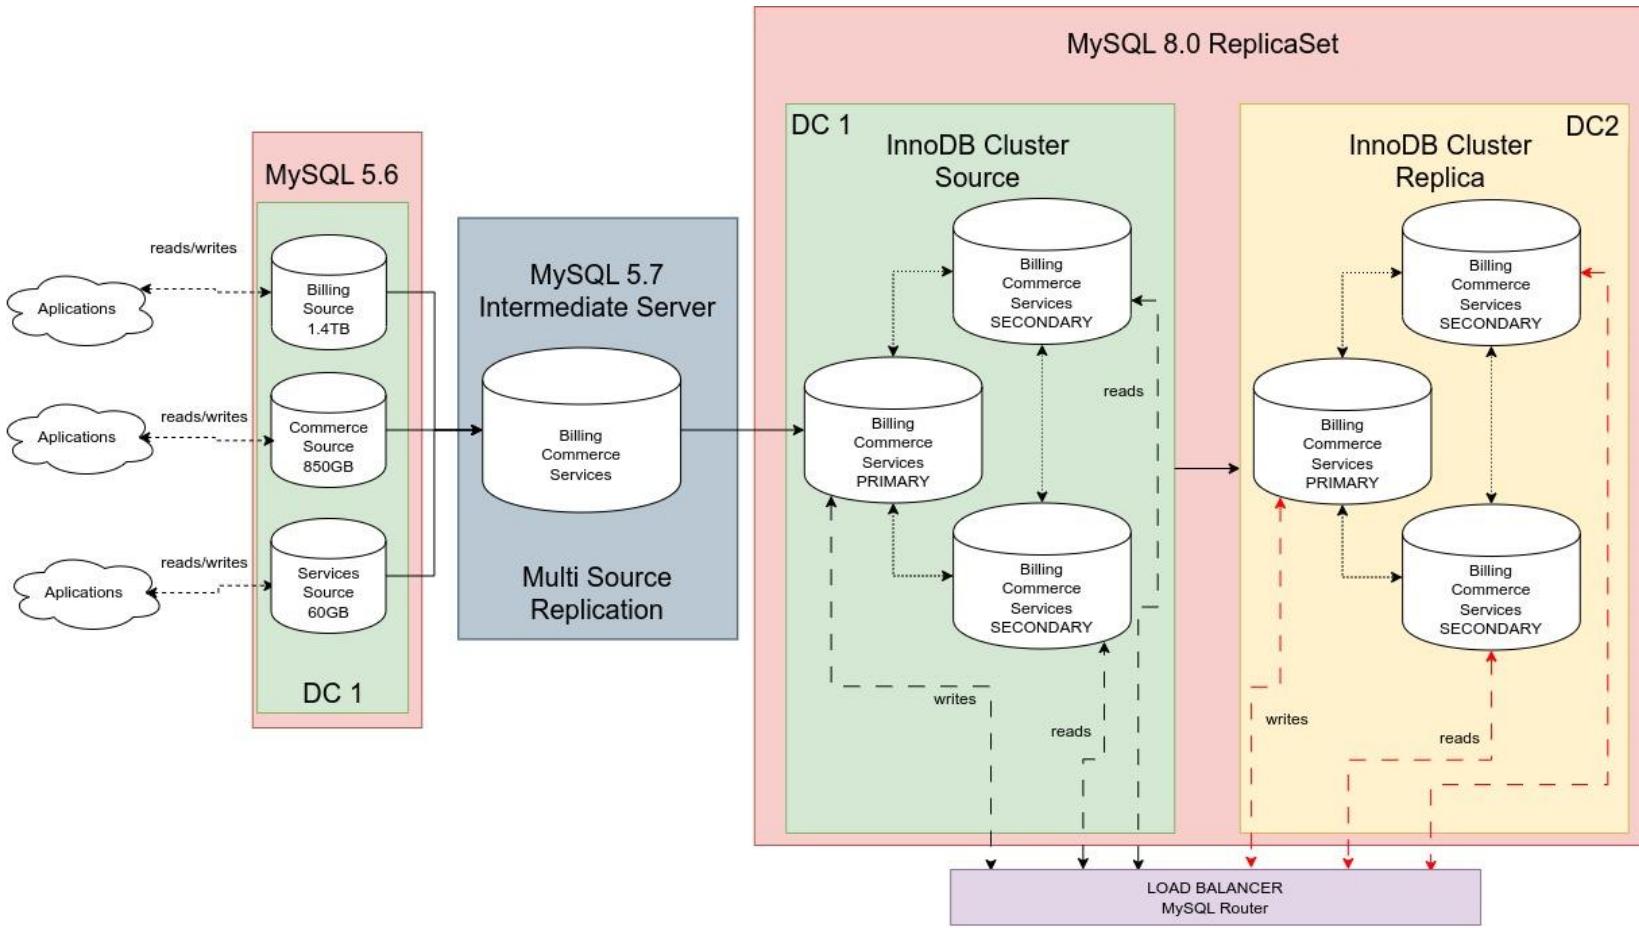
```
mysqlrouter --bootstrap root@10.50.30.101 --account=clusteradm --name='GR_Router' --force  
--directory=/home/mysql_router/mysql_router --conf-use-gr-notifications  
Please enter MySQL password for root:  
# Bootstrapping MySQL Router instance at '/home/mysql_router/mysql_router'...  
  
Please enter MySQL password for clusteradm:  
Fetching Cluster Members  
trying to connect to mysql-server at 10.50.30.101:3306  
- Creating account(s) (only those that are needed, if any)  
- Verifying account (using it to run SQL queries that would be run by Router)  
- Storing account in keyring  
- Adjusting permissions of generated files  
- Creating configuration /home/mysql_router/mysql_router/mysqlrouter.conf  
  
# MySQL Router 'GR_Router' configured for the InnoDB Cluster 'gr_MULTI_DC'  
  
After this MySQL Router has been started with the generated configuration  
$ mysqlrouter -c /home/mysql_router/mysql_router/mysqlrouter.conf  
  
InnoDB Cluster 'gr_MULTI_DC' can be reached by connecting to:  
## MySQL Classic protocol  
- Read/Write Connections: localhost:6446  
- Read/Only Connections: localhost:6447  
  
## MySQL X protocol  
- Read/Write Connections: localhost:6448  
- Read/Only Connections: localhost:6449
```

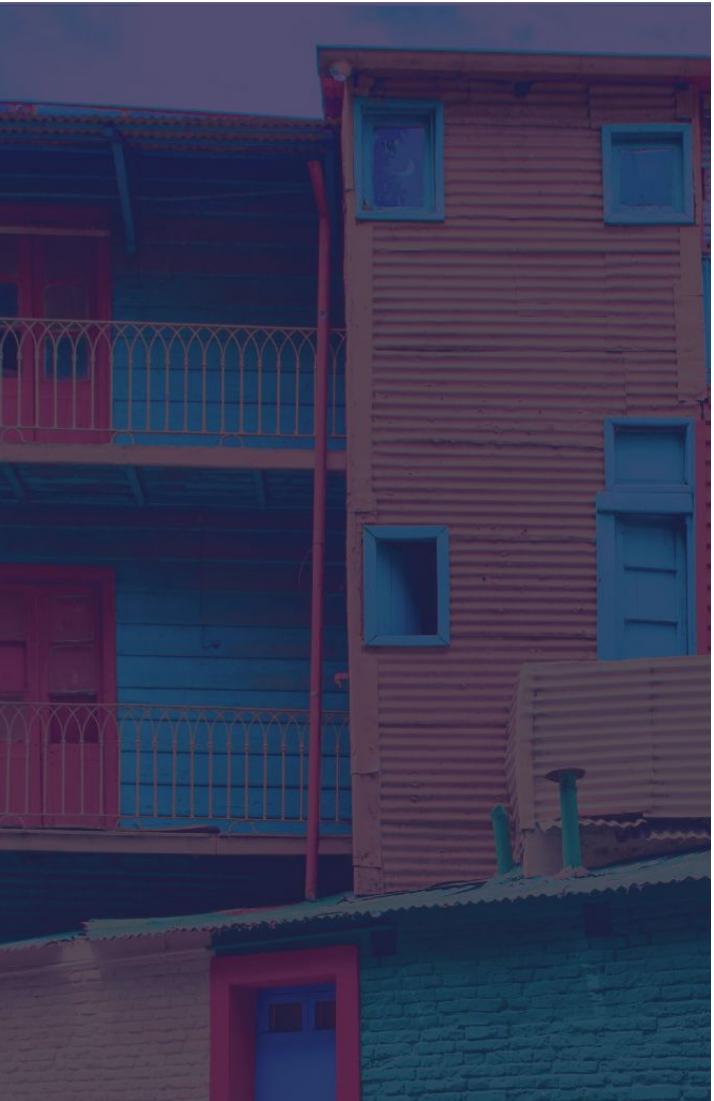
Setear replicación de 5.7 a 8.0

```
CHANGE REPLICATION SOURCE TO source_host='10.50.30.250', source_user='replication',
source_password='r3pl1c4t10n', source_log_file='mysql-bin.000181', source_log_pos=185086690,
ASSIGN_GTIDS_TO_ANONYMOUS_TRANSACTIONS='LOCAL' for channel 'production';

START REPLICA for channel 'production';
```

Escenario previo a la migración





Consideraciones finales



Podemos mejorar este setup?

Por supuesto que podemos!



Podemos mejorar este setup?

Por supuesto que podemos!

Agregar un 2do MySQL router



Podemos mejorar este setup?

Por supuesto que podemos!

Agregar un 2do MySQL router

Usar ProxySQL para balanceo de carga



Podemos mejorar este setup?

Por supuesto que podemos!

Agregar un 2do MySQL router

Usar ProxySQL para balanceo de carga

Incrementar redundancia creando un cluster de ProxySQL



Podemos mejorar este setup?

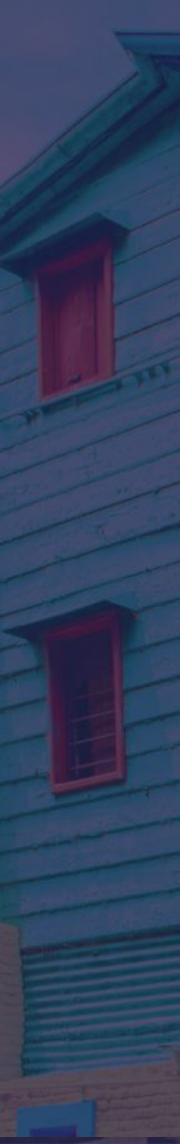
Por supuesto que podemos!

Agregar un 2do MySQL router

Usar ProxySQL para balanceo de carga

Incrementar redundancia creando un cluster de ProxySQL

Agregar una VIP por encima de la capa de conexión y administrarla con keepalived



Podemos mejorar este setup?

Por supuesto que podemos!

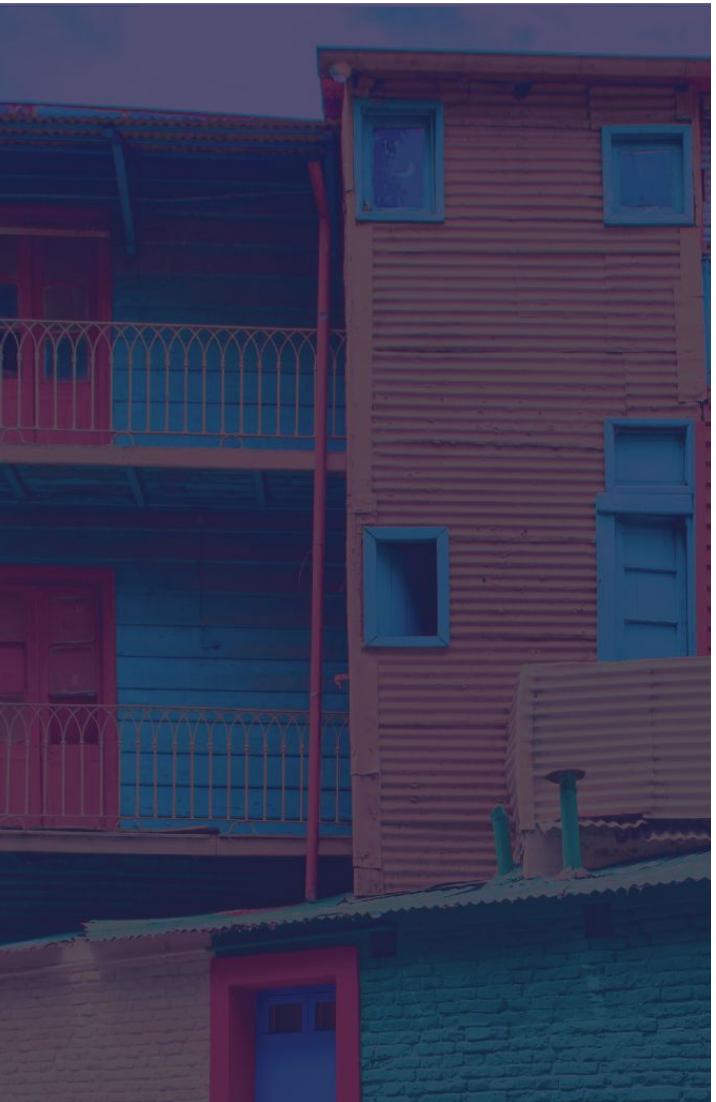
Agregar un 2do MySQL router

Usar ProxySQL para balanceo de carga

Incrementar redundancia creando un cluster de ProxySQL

Agregar una VIP por encima de la capa de conexión y administrarla con keepalived

Sean creativos!



Preguntas?





Gracias!